

Faiq Raza and Matt Robinson

## Part I Processes

The basic idea behind implementing this algorithm with processes and shared memory was to create a single contiguous array and index into it as if it was a two-dimensional array. If we would normally index with  $g[i][j]$ , then we could index into the single array with  $g[j*n + i]$ , where  $g$  was the graph,  $n$  was the number of vertices (making the matrix  $n^2$  in size), and  $i$  and  $j$  were the current nodes we were working on. The work was then split up over multiple processes by simply dividing up the rows (or columns depending on how you looked at the graph).

Here were running times for the first test case in milliseconds:

689.058899  
686.071899  
693.277466  
717.533081  
715.632202  
692.510193  
674.961792  
699.416260  
676.624695  
695.861877

## Part II Threads

We implemented the threaded part of the assignment with a different algorithm than our processed implementation. The way we set this algorithm up is figuring out the maximum work size each thread is going to be doing. We do this by doing the following calculation

$\text{max\_work\_size} = \text{vertices}/\text{threads}$

if there was a remainder the  $\text{max\_work\_size}$  was incremented by 1.

We then do a similar operation as we did in the processes implementation. Each thread gets a  $\text{work\_size}$  amount of rows to operate on. After all the threads are finished, the  $k$  is then incremented.

Some times that we got to calculate the transitive closure on test case 1 from the homework file are

132.979980  
177.298889  
175.863922  
154.878830  
465.796631  
151.326019

138.648697  
212.632629  
206.949051  
157.492676  
199.921997  
238.268265  
167.819992

There was one anomaly where the amount of time skyrocketed to 465 milliseconds